

Guide: Environment Setup for Huawei Cloud Development



Huawei Technologies CO., LTD

v3.1.0 — September 10, 2026 14:32

Contents

1 Provision a Flexus X Instance on Huawei Cloud	1
1.1 Create the security group (sg-xgate).....	1
1.2 Create the Flexus X instance	2
1.3 Verify the instance.....	3
2 Configure Flexus X for Remote Access.....	4
2.1 Set the root password.....	4
2.2 Change the SSH port to 4444.....	5
2.3 Verify the configuration.....	7
3 Install ncat and SSH Connection Topology	9
3.1 How the connection works	9
3.2 Why ncat	10
3.3 Install ncat on Windows	10
4 Install VS Code and Connect via Remote-SSH	12
4.1 Install VS Code.....	12
4.2 Install the Remote-SSH extension	12
4.3 Configure the SSH config file	12
4.4 What each ncat flag does	13
4.5 Connect to the Flexus X via VS Code.....	14
4.6 Troubleshooting	14
5 Install the MaaS Coding Gateway	16
5.1 Overview	16
5.2 Architecture	16
5.3 Run the bootstrap installer	17
5.4 Verify the installation.....	18
6 Install the Huawei Document Templates Project	19
6.1 Clone and set up the project	19
6.2 Compile the sample guides.....	20
6.3 Create your first guide with the skill.....	20
6.4 Write content using template commands.....	21
7 Clean Up.....	23
8 Changelog	24

1 Provision a Flexus X Instance on Huawei Cloud

General Objective: Provision a Flexus X instance on Huawei Cloud that will host a complete coding stack environment: the [oh-my-coding-maas-gateway](#) (a LiteLLM-based proxy with load balancing, virtual keys, and Grafana observability), governance tooling, and multiple coding agents — with Open-Code as the primary one. By the end of this chapter you will have a running Flexus X instance with a public IP accessible over SSH.

Prerequisites:

- Active Huawei Cloud account with IAM credentials.
- Permission to create security groups, Flexus X instances, and EIPs.
- A key pair created in the Console, or permission to create one.
- A web browser (Chrome, Firefox, or Edge).

Info

This guide assumes the `sa-brazil-1` region. Adjust the region if your account is registered elsewhere.

1.1 Create the security group (sg-xgate)

Objective: Create a security group that allows SSH access on port 4444 from specific Huawei internal IPs only.

Info

We use port 4444 instead of the standard SSH port 22 because the Huawei corporate proxy blocks outbound connections on port 22. Port 4444 is allowed through the proxy and works reliably for SSH.

Step by step:

1. In the Console, go to **VPC** → **Security Groups** and click **Create Security Group**.
2. Name it `sg-xgate` and add the following inbound rules:

```
Rule 1:
  Protocol:    TCP
  Port:        4444
  Source:      119.8.89.193/32
  Description: SSH for Huawei internal user 1
```

```
Rule 2:
```



```
Protocol:    TCP
Port:       4444
Source:     119.8.89.3/32
Description: SSH for Huawei internal user 2

Rule 3:
Protocol:    TCP
Port:       4444
Source:     119.8.193.3/32
Description: SSH for Huawei internal user 3
```

3. Click **OK** to create the security group.

Tip

Each /32 rule restricts access to exactly one IP address. Only the three listed Huawei internal IPs can reach the instance on port 4444 — no other inbound traffic is allowed.

1.2 Create the Flexus X instance

Objective: Launch a single Flexus X instance using the Console wizard, using the security group created above.

Step by step:

1. Log in to the [Huawei Cloud Console](#) and select the `sa-brazil-1` region in the top-right corner.
2. Open the service catalog and select **Flexus X** (under **Compute**).
3. Click **Create Flexus X** to open the creation wizard.
4. Under **Billing Mode**, select **Pay-per-use**.

Tip

Pay-per-use bills by the hour with no long-term commitment. You only pay for the time the instance is running. This is ideal for development and testing — stop the instance when not in use to reduce costs.

5. Set the compute resources using the slider: drag to **4 vCPU** and **16 GB RAM**.

Info

Flexus X uses a slider for vCPU and memory instead of fixed flavor names. The resulting flavor ID is `x1.4u.16g`, but you select it by dragging the vCPU and RAM controls — there is no flavor dropdown.

6. Fill in the remaining basic configuration:

```
AZ:          Any (choose the availability zone you prefer)
Image:       Ubuntu 24.04 Server 64bit
Disk:        100 GB General Purpose SSD
VPC:         vpc-default
Subnet:      subnet-default
Security Group: sg-xgate (port 4444, Huawei internal IPs only)
```

Important

The EIP is auto-assigned and configured to be released when the instance is deleted. If you need a fixed IP that persists across stop/start cycles, allocate a dedicated EIP and bind it manually after creation.

- Under **EIP** settings, configure the public IP:

EIP:	Auto-assign
Billing:	Pay by traffic
Bandwidth:	1000 Mbit/s
Release:	Yes – release EIP when instance is deleted

- Under **Cloud Eye**, enable monitoring (recommended).

Info

Cloud Eye is a monitoring service that provides functions such as real-time monitoring, alarm reporting, resource grouping, and website monitoring. Enabling it gives you visibility into instance health, CPU, memory, and network metrics out of the box.

- Set the instance name to `f1exusx-dev` (or any name you prefer).
- Under **Login Configuration**, select **Key pair** as the authentication method. If you do not yet have a key pair, click **Create Key Pair** to generate one inside Huawei Cloud and download the `.pem` file.

Info

The key pair is the primary access method. Later, you can set a root password on the instance (via the Console or after first login with the key) to have a second way to access the instance.

- Click **Create Now** and wait for the status to change to **Running**.

1.3 Verify the instance

On the instance details page, copy the **EIP** field. This is the address you will use to connect in the next chapters.

Step	What to check	Expected result
Security group	<code>sg-xgate</code> with TCP 4444 rules	3 inbound rules
Billing mode	Pay-per-use selected	Hourly billing
Compute	4 vCPU, 16 GB RAM via slider	Flavor x1.4u.16g
Image	Ubuntu 24.04 Server 64bit	Selected
EIP	Auto-assigned	Public IP available
Cloud Eye	Monitoring enabled	Metrics visible
Instance name	<code>f1exusx-dev</code>	Set
Key pair	Created and downloaded	<code>.pem</code> file saved
Instance	Status is Running	EIP copied

Table 1: Chapter 1 verification checklist.

2

Configure Flexus X for Remote Access

General Objective: Prepare the Flexus X instance for remote access by setting a root password via the Console and changing the SSH port from 22 to 4444. This is required because the Huawei corporate proxy blocks outbound connections on port 22.

Prerequisites:

- Chapter 1 completed — Flexus X instance running with a public IP.
 - Access to the Huawei Cloud Console.
-

2.1 Set the root password

Objective: Use the Console to reset the root password of the Flexus X instance.

Step by step:

1. In the Console, navigate to **Flexus X** and click on the instance `flexusx-dev`.
2. In the instance details page, click **More** → **Reset Password**.

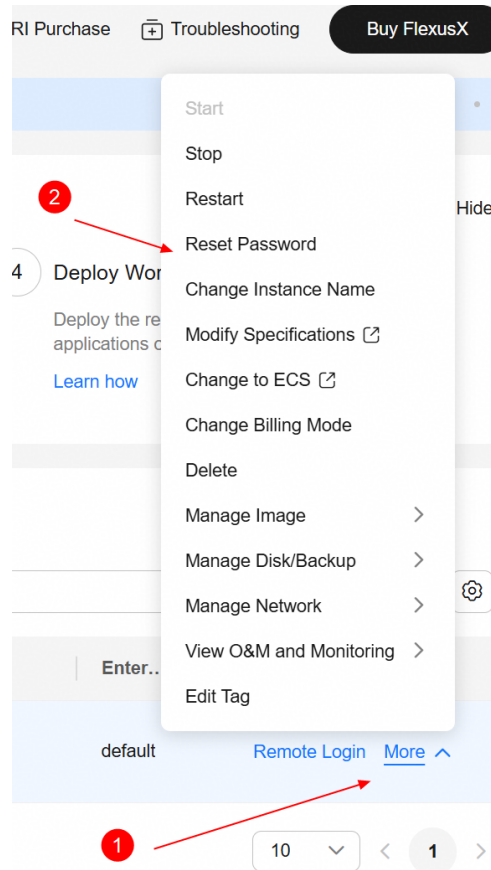


Figure 1: Reset Password option in the Flexus X details page.

3. Enter a strong root password and confirm. Click **OK**.
4. The instance will restart automatically. Wait for the status to return to **Running**.

Important

Choose a strong password (12+ characters, mixed case, digits, symbols). This password enables remote login directly from the Console and serves as a second access method alongside the SSH key pair.

2.2 Change the SSH port to 4444

Objective: Log in via the Console remote login, edit `sshd_config`, and change the SSH port from 22 to 4444.

Step by step:

1. In the instance details page, click **Remote Login**. A browser-based terminal opens.



Figure 2: Remote Login button in the Flexus X details page.

2. Log in as root using the password set in the previous subsection.
3. Open the SSH daemon configuration file:

```
sudo vim /etc/ssh/sshd_config
```

4. Find the line Port 22 (or #Port 22) and change it to:

```
Port 4444
```

5. Save and exit vim (:wq).
6. Reboot the instance from the Console to apply the change. The SSH daemon will start automatically on port 4444 after reboot.



Figure 3: Restart the Flexus X instance from the Console.

Important

The reboot is **mandatory** — not optional. Until the instance is rebooted, SSH continues listening on port 22 and the proxy CONNECT to port 4444 will fail with Ncat: Error reading proxy response Status-Line. If you skip this step, all subsequent SSH connections through the proxy will fail.

Info

After this change, SSH connections must use Port 4444 explicitly. The security group sg-xgate (created in Chapter 1) already allows inbound TCP 4444 from the authorized Huawei internal IPs.

2.3 Verify the configuration

Use the table below to confirm all steps were completed successfully before proceeding to the next chapter:



Step	What to check	Expected result
Security group	<code>sg-xgate</code> exists with TCP 4444 rules	3 inbound rules
Flexus X running	Instance <code>flexusx-dev</code> status	Running
Root password	Remote Login works with password	Shell prompt
SSH port	<code>sshd_config</code> has Port 4444	Config saved
SSH restart	Instance rebooted from Console	Service active

Table 2: Chapter 2 verification checklist.

3

Install ncat and SSH Connection Topology

General Objective: Install ncat (from Nmap) on your local Windows machine and understand the SSH connection topology through the Huawei corporate proxy (xgate). ncat serves as the SSH ProxyCommand, tunneling SSH through the HTTP proxy via the CONNECT method.

Prerequisites:

- Chapter 2 completed — Flexus X configured with root password and SSH on port 4444.
- The private key file (.pem) downloaded from Huawei Cloud.

Info

Premises:

- The proxy is **xgate** (Huawei corporate proxy) at `proxy.huawei.com:8080`.
- SSH runs on port 4444 (configured in Chapter 2).
- Login as root using the key pair from Chapter 1.
- ncat (from Nmap) is used as the ProxyCommand to tunnel SSH through the HTTP proxy via the CONNECT method.
- Proxy authentication is transparent — it is handled by the corporate VPN, not by credentials in the SSH config.

3.1 How the connection works

When you specify a ProxyCommand in the SSH config, SSH launches ncat as a child process. ncat performs the HTTP CONNECT handshake with the proxy:

```
ncat -> proxy:    CONNECT <EIP>:4444 HTTP/1.1
                  Host: <EIP>:4444

proxy -> ncat:    HTTP/1.1 200 Connection Established

[Both sides now send raw TCP bytes -- SSH negotiation begins over the tunnel]
```

After the 200 Connection Established response, the proxy becomes a blind TCP relay — it neither inspects nor modifies the SSH traffic. SSH key exchange, encryption, and authentication all happen inside this tunnel as if it were a direct TCP connection.

Info

The corporate proxy blocks CONNECT to port 22 (standard SSH) but allows it to port 4444. Enter-

prise HTTP proxies implement a CONNECT port allowlist: port 443 (HTTPS) is always allowed, port 22 is usually blocked, and port 4444 is often allowed because it is commonly used for HTTPS-alt. This is why we configured SSH on port 4444 in Chapter 2.

3.2 Why ncat

ncat is the chosen proxy tunnel tool for several reasons:

- **Full HTTP CONNECT support** — ncat implements the CONNECT method natively via `--proxy-type http`.
- **Actively maintained** — part of the Nmap project, with regular releases and bug fixes.
- **IPv4 forcing** — the `-4` flag prevents failed IPv6 attempts through the proxy.
- **Connect timeout** — the `-w` flag prevents indefinite hangs if the proxy is unreachable.
- **Proxy auth support** — `--proxy-auth` and the `NCAT_PROXY_AUTH` environment variable, if explicit credentials are ever needed.

Info

`connect.exe` (from Git for Windows) is a lighter alternative that ships with Git, but it lacks the `-4` (force IPv4) and `-w` (connect timeout) flags, which are important for reliable proxied connections. ncat provides better control over the tunnel behavior.

3.3 Install ncat on Windows

Objective: Install Nmap, which provides `ncat.exe` — the proxy tunnel tool used by SSH ProxyCommand.

Step by step:

1. Download the Nmap installer from nmap.org/download.html. Use the **Latest stable release self-installer** (`nmap-7.991-setup.exe`, tested and verified).
2. Run the installer (right-click **Run as Administrator**).
3. During installation, uncheck **Npcap** — it is a kernel driver for packet capture and is not needed for proxy tunneling. Also uncheck **Zenmap**, **Nping**, and **Ndiff** for a minimal install.



Figure 4: Nmap installer component selection — uncheck Npcap, Zenmap, Nping, and Ndiff.

Tip

Skipping Npcap avoids a kernel driver install and potential reboot. ncat in proxy mode uses standard Winsock2 TCP sockets and does not need Npcap at all.

4. Choose the default install path: C:\Program Files (x86)\Nmap.
5. After installation, verify that ncat.exe is available:

```
"C:\Program Files (x86)\Nmap\ncat.exe" --version
```

6. If the command shows the version string, the installation is correct.

Important

Nmap is commonly flagged by antivirus software due to port-scanning heuristics. Windows Defender may quarantine ncat.exe on download or block execution. Add an exclusion for the Nmap folder and the ncat.exe process in **Windows Security** → **Virus & threat protection** → **Exclusions**. If ncat.exe was already quarantined, restore it from Protection history before adding the exclusion.

Step	What to check	Expected result
Nmap download	nmap-7.991-setup.exe saved	Installer ready
Components	Npcap, Zenmap, Nping, Ndiff unchecked	Minimal install
Install path	C:\Program Files (x86)\Nmap	Default path
ncat.exe	ncat --version shows version	Installed correctly
Antivirus	Exclusion added for Nmap folder	ncat.exe not blocked

Table 3: Chapter 3 verification checklist.

4 Install VS Code and Connect via Remote-SSH

General Objective: Install VS Code from scratch on your local Windows machine, install the Remote-SSH extension, configure the SSH config file using VS Code as the editor, and connect to the Flexus X instance.

Prerequisites:

- Chapter 3 completed — ncat installed and proxy tunnel understood.
 - The private key file (.pem) downloaded from Huawei Cloud.
-

4.1 Install VS Code

Step by step:

1. Download VS Code from code.visualstudio.com/download. Choose the **Windows** installer (User Installer, 64-bit).
2. Run the installer with the default settings.
3. After installation, verify VS Code is available:

```
code --version
```

4.2 Install the Remote-SSH extension

Step by step:

1. Open VS Code on your local Windows machine.
2. Install the **Remote - SSH** extension (publisher: Microsoft):

```
code --install-extension ms-vscode-remote.remote-ssh
```

4.3 Configure the SSH config file

Objective: Create an SSH config entry that tunnels through the Huawei proxy using ncat from Nmap. Use VS Code as the editor.

Step by step:

1. Open or create the SSH config file using VS Code:

```
code %USERPROFILE%\.ssh\config
```

2. Add the following configuration. Replace these placeholders with your values:

- <host-name> — a name of your choice (e.g., flexusx-dev)
- <EIP> — the Flexus X public IP copied from Chapter 1
- <user> — your Windows username (e.g., w00946241)
- <key-file> — the name of your .pem key file downloaded from Huawei Cloud (e.g., KeyPair-opencode)

```
Host <host-name>
  HostName <EIP>
  User root
  Port 4444
  IdentityFile "C:\Users\<user>\.ssh\<key-file>.pem"
  ProxyCommand "C:\Program Files (x86)\Nmap\ncat.exe" -4 -w 10 --proxy
  proxy.huawei.com:8080 --proxy-type http %h %p
  ServerAliveInterval 15
  ServerAliveCountMax 4
  ConnectTimeout 10
  TCPKeepAlive yes
  StrictHostKeyChecking accept-new
```

Info

The IdentityFile path has two placeholders: <user> is your Windows username (the folder under C:\Users\), and <key-file> is the name of the .pem file downloaded from Huawei Cloud in Chapter 1. For example: C:\Users\w00946241\.ssh\KeyPair-opencode.pem.

4.4 What each ncat flag does

Flag	Purpose
-4	Force IPv4 (avoids failed IPv6 attempts through the proxy)
-w 10	10-second connect timeout (prevents indefinite hang)
--proxy	Proxy address and port (proxy.huawei.com:8080)
--proxy-type http	Use HTTP CONNECT method
%h	Substituted by SSH with the HostName (the EIP)
%p	Substituted by SSH with the Port (4444)

Table 4: ncat ProxyCommand flags.

Tip

The ServerAliveInterval 15 and ServerAliveCountMax 4 settings send a keepalive every 15 seconds and disconnect after 4 missed keepalives (60 seconds dead detection). This is critical for proxied connections, which may drop idle sessions after 5–30 minutes.

4.5 Connect to the Flexus X via VS Code

Step by step:

1. Press F1 to open the Command Palette, type **Remote-SSH: Connect to Host...**, and press Enter.




Figure 5: VS Code Command Palette — select Remote-SSH: Connect to Host...

2. Select your host name (e.g., flexusx-dev) from the list of configured hosts.
3. If prompted to select the platform type, choose **Linux**. This tells VS Code which server binary to download and avoids auto-detection delays.

Info

VS Code only asks for the platform type the first time you connect to a new host. The selection is saved in your local settings (`remote.SSH.remotePlatform`) and reused on subsequent connections.

4. VS Code opens a new window and begins connecting. On the first connection, it downloads and installs the VS Code Server on the remote instance — this takes about 1–2 minutes. Subsequent connections are faster.
5. Once connected, click **Open Folder** and select `/home` (or your working directory).
6. Verify the connection by opening a terminal in VS Code (`Ctrl+``) and running:

```
uname -a
df -h
```

Info

The first connection installs the VS Code Server on the remote instance, which requires outbound internet access on port 443 (HTTPS). The Flexus X instance has an EIP, so it typically has outbound internet access. If the connection hangs during server installation, verify with `curl -sI https://update.code.visualstudio.com` from the instance.

4.6 Troubleshooting

If the connection fails, check these common issues:

- **Ncat: Error reading proxy response** — SSH was not restarted after the port change (Chapter 2). **Fix:** reboot the instance from the Console.
- **Ncat: Error reading proxy response** — proxy CONNECT permission missing for this IP. **Fix:** request IT permission for the instance EIP on port 4444.
- **Connection timed out** — proxy unreachable or CONNECT blocked. **Fix:** verify proxy address; test ncat with the `-v` flag.
- **UNPROTECTED PRIVATE KEY FILE** — key permissions too open. **Fix:** run `icaccls` to restrict the key file to your user only.
- **VS Code hangs on first connect** — instance cannot reach internet on port 443. **Fix:** verify `curl -sI https://update.code.visualstudio.com` from the instance.
- **\$PLATFORM is undefined** — SSH connection failed before server install. **Fix:** check the Remote-SSH output log; fix the underlying SSH error first.

Info

To debug the proxy connection, run ncat directly with verbose output:

```
"C:\Program Files (x86)\Nmap\ncat.exe" -v --proxy proxy.huawei.com:8080
--proxy-type http <EIP> 4444
```

This shows the CONNECT request and the proxy's raw response (or connection close). For SSH-level debugging, run `ssh -vvv <host-name> "echo ok"`.

Step	What to check	Expected result
VS Code	<code>code --version</code> works	Installed
Remote-SSH	Extension installed	Active
SSH config	Host entry with ncat ProxyCommand	Config saved
Connection	Remote-SSH connects to host	Flexus X accessible
Terminal	<code>uname -a</code> and <code>df -h</code>	Output verified

Table 5: Chapter 4 verification checklist.

5

Install the MaaS Coding Gateway

General Objective: Install the oh-my-coding-maas-gateway on the Flexus X instance. This is a self-hosted LiteLLM proxy that routes Huawei ModelArts MaaS models to coding tools with load balancing, virtual keys, dual-format API endpoints, and Prometheus + Grafana observability.

Prerequisites:

- Chapter 4 completed — VS Code connected to the Flexus X via Remote-SSH.
 - Docker and Docker Compose available (the installer handles this).
 - A Huawei ModelArts MaaS API key.
-

5.1 Overview

The oh-my-coding-maas-gateway is a self-hosted proxy gateway that sits between your coding tools and Huawei ModelArts MaaS (Model-as-a-Service) models. It provides:

- **Load balancing** — distributes requests across multiple MaaS API keys (N keys → N deployments per model per format).
- **Virtual key management** — each coding tool gets its own virtual key with unlimited budget, isolated from the real MaaS keys.
- **Dual-format API endpoints** — OpenAI Chat Completions (/v1/chat/completions) and Anthropic Messages (/v1/messages), so tools using different API formats can all reach the same MaaS models.
- **Budget tracking** — per-key spend tracking via LiteLLM + PostgreSQL, with cost-per-token configured for all models.
- **Full observability** — Prometheus metrics + Grafana 39-panel dashboard (latency, errors, throughput, tokens, cache, cost).

Info

The gateway repository is at [wallacelw/oh-my-coding-maas-gateway](https://github.com/wallacelw/oh-my-coding-maas-gateway). It supports 4 models: glm-5.2 (1M context), glm-5.1, deepseek-v4-pro, and deepseek-v4-flash (384K max output). This guide was tested with v1.10.0; the installation process is the same for newer versions.

5.2 Architecture

The gateway deploys 4 Docker containers via Docker Compose:

Tools	LiteLLM (:4000)	Huawei MaaS
opencode	--> /v1/chat/completions	--> openai/ --> MaaS OpenAI endpoint
Codex CLI	--> /v1/responses	--> openai/ --> MaaS OpenAI endpoint
Claude Code endpoint	--> /v1/messages	--> anthropic/ --> MaaS Anthropic endpoint
Pi agent	--> /v1/chat/completions	--> openai/ --> MaaS OpenAI endpoint
LiteLLM: load-balances across N MaaS keys . PostgreSQL (:5432)		
Observability: LiteLLM --/metrics--> Prometheus (:9090) --> Grafana (:3000)		

Service	Port	Purpose
LiteLLM Proxy	4000	API gateway (OpenAI + Anthropic format)
PostgreSQL	5432	LiteLLM database (keys, spend, config)
Prometheus	9090	Metrics storage and querying
Grafana	3000	Dashboard visualization (39 panels)

Table 6: Docker services and ports.

Info

All ports bind to 127.0.0.1 (localhost only) by default. To access Grafana and the LiteLLM Admin UI from your local machine, use SSH port forwarding in VS Code or add port forwards in the Remote-SSH connection.

Each model has two deployments — one OpenAI, one Anthropic — so all tools can use the same MaaS models regardless of their API format. The `claude-` prefix (e.g., `claude-glm-5.2`) prevents routing conflicts in LiteLLM.

5.3 Run the bootstrap installer

Objective: Install the gateway using the one-command bootstrap script.

Step by step:

1. In the VS Code remote terminal, run the bootstrap script:

```
curl -fsSL https://raw.githubusercontent.com/wallacelw/oh-my-coding-maas-gateway/main/scripts/bootstrap.sh | bash
```

2. The installer will:

- Prompt for your Huawei MaaS API key.
- Show an interactive menu to select which coding tools to install (opencode, Codex CLI, Claude Code, Pi agent).
- Auto-install prerequisites (Docker, Node.js, etc.).
- Deploy the proxy + observability stack (LiteLLM, PostgreSQL, Prometheus, Grafana).
- Configure each tool with its own virtual key.
- Run 120 validation checks and offer to install the companion skill.

Important

Estimated time: ~5 min for a fresh install, ~2 min for an upgrade. If you are behind a proxy, ensure `http_proxy` and `https_proxy` environment variables are set before running the script.

Tip

For non-interactive installs (CI/automation), pass the API key directly:

```
curl -fsSL https://raw.githubusercontent.com/wallacelw/oh-my-coding-maas-gateway/main/scripts/bootstrap.sh | bash -s -- -y --api-key=sk-xxxx
```

5.4 Verify the installation

Step by step:

1. Check that the proxy is running:

```
curl http://localhost:4000/health
```

2. Open the Grafana dashboard at `http://localhost:3000` to monitor usage and metrics (39 panels across 7 sections).
3. Open the LiteLLM Admin UI at `http://localhost:4000/ui` to manage virtual keys and model routing.
4. Launch your coding tool:

```
opencode # or: codex or: claude --bare or: pi
```

Tip

To upgrade the gateway later, re-run the same bootstrap command. It detects existing installs, compares versions, and pulls updates — all secrets and data are preserved.

Step	What to check	Expected result
Bootstrap	Installer completed	All services deployed
Proxy health	<code>curl localhost:4000/health</code>	Healthy
Grafana	<code>localhost:3000</code>	Dashboard accessible
LiteLLM UI	<code>localhost:4000/ui</code>	Admin UI accessible
Coding tool	Launched and connected	Gateway routing

Table 7: Chapter 5 verification checklist.

6 Install the Huawei Document Templates Project

General Objective: Clone the Huawei document templates repository on the Flexus X instance, run the setup script, and verify that the LaTeX toolchain and the coding agent skill are ready for use. The script assumes a fresh Ubuntu install — it installs all dependencies from scratch and configures VS Code at the user level.

Prerequisites:

- Chapter 5 completed — MaaS gateway installed and running.
- Git installed (`sudo apt install git`).
- The repository URL for the Huawei document templates project.
- A fresh Ubuntu 24.04 instance (no prior LaTeX installation required).

Info

The `install.sh` script is idempotent — it is safe to run multiple times. It merges VS Code settings rather than overwriting them, and `apt-get install` skips packages that are already present.

6.1 Clone and set up the project

Step by step:

1. Clone the repository:

```
cd ~
git clone https://github.com/wallacelw/huawei-doc-templates.git
cd huawei-doc-templates
```

2. Run the setup script (installs XeLaTeX, latexmk, fonts, the coding agent skill, and configures VS Code LaTeX Workshop at the user level):

```
./install.sh
```

Important

The script runs `sudo apt-get install` for TeX Live packages. If you do not have sudo access, ask your administrator to install the packages listed in `install.sh`.

3. Verify the LaTeX toolchain:

```
xelatex --version
latexmk --version
```

6.2 Compile the sample guides

The project includes a self-documenting Makefile. To compile all samples (guide + technical, Portuguese + English):

Step by step:

1. From the repository root, run:

```
make samples
```

2. Verify the PDFs were generated:

```
ls -l examples/guide/pt/main.pdf
ls -l examples/guide/en/main.pdf
ls -l examples/technical/pt/main.pdf
ls -l examples/technical/en/main.pdf
```

Tip

To compile a single project, use `make project DIR=examples/guide/pt` or run `latexmk main.tex` from the project's `src/` directory. Run `make` with no arguments to see all available targets.

6.3 Create your first guide with the skill

Step by step:

1. In the VS Code remote terminal, launch your coding agent:

```
cd ~/huawei-doc-templates
opencode # or: codex, claude --bare, pi
```

2. Run the guide skill:

```
/skill huawei-template-guide
```

3. The skill will ask for:

- **Title** — e.g. “Provisioning an OBS Bucket”
- **Language** — Portuguese or English
- **Project name** — used as the folder name
- **Location** — where to create the project folder

4. The skill generates the `.tex` file, creates a `.latexmkrc` with the correct `TEXINPUTS`, compiles the document, and reports the page count.

Ready

Info

The skill is auto-discovered via `opencode.json` which registers `templates/` as a skill discovery path. No manual skill installation is needed when working inside the repo.

6.4 Write content using template commands

The template provides commands for common guide elements:

Headings (automatic numbering, H1 starts on a new page):

```
\section{Chapter Title}          % H1: giant number + title + rule
\subsection{Section Title}      % H2
\subsubsection{Subsection Title} % H3
```

Objectives block (closes with a horizontal rule):

```
\begin{objectives}
  \generalobjective{...}
  \prerequisites
  \begin{itemize}
    \item ...
  \end{itemize}
\end{objectives}
```

Step-by-step and objectives (used inside subsections):

```
\objective{...}
\stepbystep
\begin{enumerate}
  \item ...
\end{enumerate}
```

Callout boxes (environment syntax):

- `warning` — amber warning box
- `tip` — green tip box
- `infobox` — blue info box

Code blocks and inline commands:

- `\begin{code}` — code block (verbatim, no escaping needed inside)
- `\begin{code}[bash]` — same, with a language hint
- `\inlinecode{...}` — inline monospace code

Images, links, and other inline commands:

```
\image{assets/screenshot.png}          % centered image
\imagecap{assets/screenshot.png}{caption} % image with numbered caption
\note{Italic observation.}           % note paragraph
\weblink{https://...}{link text}        % blue clickable link
\menu{VPC, Security Groups}             % menu path: VPC -> Security Groups
\badge{New}                             % inline red label
```

Info

The project also includes a technical template for technical reports. Use `/skill huawei-template-technical` to create technical reports with a 5-section structure (problem, root cause analysis, root cause, trigger condition, workaround). See `templates/technical/SKILL.md` for the full command reference.

Step	What to check	Expected result
Repository	Cloned to <code>~/huawei-doc-templates</code>	Files present
<code>install.sh</code>	Completed without errors	Dependencies installed
XeLaTeX	<code>xelatex --version</code> works	Toolchain ready
latexmk	<code>latexmk --version</code> works	Build system ready
Sample PDFs	<code>make samples</code> compiles all	Guide + technical, PT + EN
Skill	<code>/skill huawei-template-guide</code>	Available

Table 8: Chapter 6 verification checklist.

7

Clean Up

General Objective: Remove the resources created in this guide to avoid ongoing charges on your Huawei Cloud account.

Step by step:

1. In the Console, navigate to **Flexus X**.
2. Select the instance `f1exusx-dev`, click **More** and choose **Delete**.
3. Confirm deletion. The EIP is released automatically (configured in Chapter 1 to release on instance deletion).
4. Navigate to **VPC** → **Security Groups**, select `sg-xgate`, and delete it.
5. Remove the SSH config entry for `f1exusx-dev` from your Windows SSH config file (`C:\Users\<user>\.ssh\co`).

Important

If you used Terraform to provision the infrastructure, run `terraform destroy` instead to remove all resources declared in the configuration.

Step	What to check	Expected result
Flexus X instance	<code>f1exusx-dev</code> deleted	No running instance
EIP	Released	No public IP charges
Security group	<code>sg-xgate</code> deleted	No SG charges
SSH config	Host entry removed	Clean config

Table 9: Chapter 7 verification checklist.

This guide was generated using the `guide.cls` template. For the full command reference, see the template's `README.md` or run `/skill huawei-template-guide` in your coding agent.

3.1.0*September 10, 2026*

- Switched from ECS to Flexus X instance (4 vCPU, 16 GB RAM via slider, flavor x1.4u.16g) — more cost-efficient and robust. Updated creation steps to match Flexus X console order: billing mode first, vCPU/RAM slider (no flavor dropdown), Cloud Eye monitoring (recommended), then instance name and key pair login. Replaced screenshots with Flexus X console images.
- Rewrote Chapter 3: switched from `connect.exe` to `ncat` (from Nmap) as the SSH ProxyCommand. Added explanation of HTTP CONNECT method, why port 4444 works, why `ncat` is chosen, `ncat` installation with antivirus exclusions, flag reference table, and improved keepalive settings.
- Restructured chapters: moved SSH config and VS Code connection to new Chapter 4 (Install VS Code and Connect via Remote-SSH). Chapters 5–7 renumbered.
- Added verification checklist tables to all 7 chapters.
- Added mandatory reboot warning in Chapter 2 (SSH port change). Added troubleshooting subsection in Chapter 4 with common SSH connection issues and `ncat` debugging commands.

3.0.1*September 10, 2026*

- Minor fixes: `test-filter.sh` auto-discovery, `build.sh` help text, Makefile aggregate descriptions, 5-section comment correction, CHANGELOG version gap note, callout-in-code fix, CI font auto-discovery, `dofile` path comment, `test-sync.sh` header.

3.0.0*September 10, 2026*

- Modular pipeline refactoring: shared Lua filter factory, shared DOCX fix logic, build system auto-discovery, unified format pipeline, GitHub Actions CI.

2.16.0*September 10, 2026*

- Updated to match template v2.16.0 — Lua filter fixes, build system updates, and technical template parity improvements.

2.13.0*2026-08-24*

- Foundation refactoring Phase 4: make DOCX `--fix` post-processing robust. Replaced all silent-skip guards in `create-reference-docx.py` with loud `RuntimeError` assertions for expected XML elements (Heading1–4, Title, VerbatimChar, Normal, docDefaults, numbering indentation). Added pandoc version pin (3.1.0–3.2.0) to catch output-structure regressions on upgrade. Extended `test-docx-fix.sh` with version check and 4 loud-failure tests (missing style → non-zero exit).

2.12.0*2026-08-24*

- Foundation refactoring Phase 3: introduce `parse_preamble()` in the Lua filter, extracting all class options (`portuguese`, `indentbody`, `notime`, `nochangelog`) and all `\set*` command values into a single

preamble state object. Fixes 6 L18 divergences at root cause: English TOC label (“Table of Contents” → “Contents”), [notime] now gates time in DOCX/HTML, [nochangelog] now suppresses changelog in non-PDF formats, \setcoverlogo honored instead of hardcoded, \setheaderlogo warns if set in non-PDF, \note renders as italic in all formats (matching PDF). Any future preamble command or class option now flows through automatically.

2.11.0

2026-08-24

- Foundation refactoring Phase 2: extend test matrix to cover all 3 output formats × 3 documents (9 cells). `round-trip.sh` now generates DOCX, counts structures in `word/document.xml`, and checks cross-format consistency (headings, images, code blocks, tables, callouts, 0 raw LaTeX across MD+DOCX+HTML). Added `test-docx-fix.sh` smoke test asserting H1 red border, heading colors, list indentation, and footer PAGE field in patched `styles.xml`.

2.10.0

2026-08-24

- Foundation refactoring Phase 1: sync cls version to project tag, fix stale doc references (L16 output count, L-range, setup-guide feature claim), add `test-sync.sh` for version consistency checking.

2.9.1

2026-08-23

- HTML and Markdown outputs now embed images as base64 data URIs. HTML uses `pandoc --embed-resources`, MD uses post-processing with `embed-images.py`. Outputs are self-contained — no external image files needed.

2.9.0

2026-08-23

- DOCX: fix content width (8504→9638tw, margins are 2cm not 3cm). Add footer page numbers. Style TOC heading (22pt bold + rule). Add hutable table width + fixed layout.

2.8.3

2026-08-23

- DOCX fix: callout tables now have explicit width (8504tw) and fixed layout. Changelog section heading now uses custom H1 format (56pt number, red border, right tab) instead of pandoc’s default.

2.8.2

2026-08-23

- DOCX H1 heading: reverted table approach to paragraph with tab stop + red bottom border. Table broke navigation pane and TOC field. Paragraph preserves Heading1 style for navigation while matching PDF layout (number left, title right, red rule below).

2.8.1

2026-08-23

- DOCX heading fix: text color reverted to black (was incorrectly changed to Huawei red in v2.8.0). H1 now uses a two-column table (number left, title right) with red bottom rule. H2-H4 use simple inline format (number + title, no right tab stop).

2.8.0

2026-08-23

- DOCX output: match PDF styling — heading colors (Huawei red), H1 number size (56pt), callout borders (left-only) and padding, changelog formatting (rules, bold version, italic date), caption style (9pt bold), badge style (8pt bold white on red), list indentation.

2.7.1

2026-08-17

- Rename project repository to all lowercase (`huawei-doc-templates`).

2.7.0	2026-08-12
• The <code>changelog</code> environment now emits its own section heading; <code>[nochangelog]</code> suppresses the heading and entries in one switch.	
2.6.1	2026-08-12
• Fix <code>hutable</code> centering and row-color fill (switched to <code>\hline</code> for <code>colortbl</code> compatibility).	
2.6.0	2026-08-12
• Switched verification table to the new <code>hutable</code> full-grid style.	
• Floats now default to <code>[H]</code> placement (in-source order).	
2.5.1	2026-08-10
• H1 section title size reduced from 27pt to 20pt.	
2.5.0	2026-08-10
• Default image size increased: width 65% → 90% of text width, height 40% → 50% of text height.	
2.4.0	2026-08-10
• Enlarged H1 section title (18pt → 27pt, 50% bigger).	
• <code>nochangelog</code> option now hides version, date, and time on the cover page.	
2.3.1	2026-08-09
• Rolled back <code>/ActualText</code> markers — broke code block rendering.	
2.3.0	2026-08-09
• Added key-value sizing options to <code>\image</code> and <code>\imagecap</code> — width and height can now be set independently.	
2.2.4	2026-08-09
• Fixed PDF copy-paste of URLs in code blocks — disabled Cascadia Code programming ligatures (<code>ca1t</code>) that corrupted <code>//</code> into <code>))</code> on copy.	
2.2.3	2026-08-09
• Fixed italic font rendering — HarmonyOS Sans and Cascadia Code have no italic variants; enabled synthetic slant (<code>AutoFakeSlant</code>).	
2.2.2	2026-08-08
• Fixed chapter reference in prerequisite (Chapter 3 → Chapter 4).	
2.2.1	2026-08-08
• Changed table and figure caption labels to black (was Huawei red).	
2.2.0	2026-08-08
• Added Huawei-branded table styling: red header bar, alternating row colors, and bordered cells.	
2.1.0	2026-08-06



- Updated ECS flavor to `ac8.xlarge.2` (4 vCPU, 8 GB RAM).
- Updated OS image to Ubuntu 24.04 Server 64bit.
- Increased disk size to 100 GB.
- Recommended pay-per-use billing mode.
- EIP auto-assigned: pay by traffic, 1000 Mbit/s, release with ECS.
- Key pair authentication (created in Huawei Cloud); root password can be set later as a second access method.
- Security group `sg-xgate` now created **before** the ECS instance.
- Replaced SSH port 22 with port 4444 (Huawei proxy blocks port 22).
- Restricted SSH access to three specific Huawei internal IPs (119.8.89.193, 119.8.89.3, 119.8.193.3).
- Added Port 4444 to SSH config examples.

2.0.0

2026-08-05

- Replaced OpenCode with `oh-my-coding-maas-gateway` as the default coding environment.
- Added changelog and versioning support to the template (`\setdocversion`, `\setdocdate`, `changelog` environment).
- Renamed `info` environment to `infobox` to avoid package name collisions.
- Fixed external file code block implementation (`\codefile`).

1.0.0

2026-08-05

- Initial version.
 - Added ECS provisioning on Huawei Cloud.
 - Added SSH with proxy configuration for VS Code Remote-SSH.
 - Replaced OpenCode with `oh-my-coding-maas-gateway` (LiteLLM proxy with load balancing, virtual keys, and Grafana observability).
 - Added Huawei document templates project setup.
 - Added changelog and versioning support to the template.
-